

# LAB REPORT: 2

## Gradient Descent and Backpropagation with Different Activation Functions

**Name:** Aayush Suthar

**Registration No.:** 23FE10CAI00275

**Course:** Deep Learning Laboratory

**Semester:** VI | **Section:** F

**Date:** 15th January 2026

---

### 1. OBJECTIVE

To implement gradient descent and backpropagation algorithms with different activation functions (Sigmoid, Tanh, ReLU) for training neural networks.

---

### 2. THEORY

#### Gradient Descent

An optimization algorithm that minimizes loss by iteratively adjusting weights in the direction of steepest descent.

**Update Rule:**  $W = W - \alpha \cdot \nabla_W L$

#### Backpropagation

Efficient algorithm for computing gradients by applying the chain rule backward through the network.

#### Activation Functions

Function	Formula	Derivative	Range
Sigmoid	$\sigma(z) = 1/(1+e^{-z})$	$\sigma(z) \cdot (1-\sigma(z))$	$(0, 1)$
Tanh	$\tanh(z) = (e^z - e^{-z})/(e^z + e^{-z})$	$1 - \tanh^2(z)$	$(-1, 1)$
ReLU	$\max(0, z)$	1 if $z>0$ else 0	$[0, \infty)$

---

### 3. CODE IMPLEMENTATION

```
python
```

```

import numpy as np
import matplotlib.pyplot as plt

np.random.seed(42)

# XOR Dataset
X = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
y = np.array([[0], [1], [1], [0]])

# Activation Functions
def sigmoid(z): return 1 / (1 + np.exp(-np.clip(z, -500, 500)))
def sigmoid_prime(z): s = sigmoid(z); return s * (1 - s)
def tanh(z): return np.tanh(z)
def tanh_prime(z): return 1 - np.tanh(z) ** 2
def relu(z): return np.maximum(0, z)
def relu_prime(z): return (z > 0).astype(float)

# Neural Network Class
class NeuralNetwork:
    def __init__(self, activation='sigmoid', lr=0.5):
        self.W1 = np.random.randn(2, 4) * 0.5
        self.b1 = np.zeros((1, 4))
        self.W2 = np.random.randn(4, 1) * 0.5
        self.b2 = np.zeros((1, 1))
        self.lr = lr
        self.losses = []

    # Select activation
    if activation == 'sigmoid':
        self.act, self.act_prime = sigmoid, sigmoid_prime
    elif activation == 'tanh':
        self.act, self.act_prime = tanh, tanh_prime
    else: # relu
        self.act, self.act_prime = relu, relu_prime

    def forward(self, X):
        self.z1 = X @ self.W1 + self.b1
        self.a1 = self.act(self.z1)
        self.z2 = self.a1 @ self.W2 + self.b2
        self.a2 = sigmoid(self.z2)
        return self.a2

    def backward(self, X, y):
        m = X.shape[0]
        dz2 = self.a2 - y
        dW2 = (1/m) * self.a1.T @ dz2

```

```

db2 = (1/m) * np.sum(dz2, axis=0, keepdims=True)

da1 = dz2 @ self.W2.T
dz1 = da1 * self.act_prime(self.z1)
dW1 = (1/m) * X.T @ dz1
db1 = (1/m) * np.sum(dz1, axis=0, keepdims=True)

return dW1, db1, dW2, db2

def train(self, X, y, epochs=10000):
    for epoch in range(epochs):
        y_pred = self.forward(X)
        loss = -np.mean(y * np.log(y_pred + 1e-15) +
                        (1 - y) * np.log(1 - y_pred + 1e-15))
        self.losses.append(loss)

        dW1, db1, dW2, db2 = self.backward(X, y)
        self.W1 -= self.lr * dW1
        self.b1 -= self.lr * db1
        self.W2 -= self.lr * dW2
        self.b2 -= self.lr * db2

        if (epoch + 1) % 2000 == 0:
            print(f"Epoch {epoch+1}, Loss: {loss:.4f}")

    def predict(self, X):
        return (self.forward(X) >= 0.5).astype(int)

# Train with different activations
activations = ['sigmoid', 'tanh', 'relu']
models = {}

for act in activations:
    print(f"\n{'='*50}\nTraining with {act.upper()}\n{'='*50}")
    models[act] = NeuralNetwork(activation=act)
    models[act].train(X, y)

# Compare Results
print("\n" + "="*60)
print("PERFORMANCE COMPARISON")
print("="*60)

for act in activations:
    pred = models[act].predict(X)
    prob = models[act].forward(X)
    acc = np.mean(pred == y) * 100
    loss = models[act].losses[-1]

```

```
print(f"\n{act.upper()}: Loss={loss:.4f}, Accuracy={acc:.0f}%")
print("Input\tTarget\tPred\tProb")
for i in range(4):
    print(f" {X[i]}\t{y[i][0]}\t{pred[i][0]}\t{prob[i][0]:.3f}")

# Visualize
plt.figure(figsize=(10, 5))
for act in activations:
    plt.plot(models[act].losses, label=act.upper())
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.title('Training Loss Comparison')
plt.legend()
plt.grid(True)
plt.show()
```

## Output:

---

---

### Training with SIGMOID

---

---

Epoch 2000, Loss: 0.1089  
Epoch 4000, Loss: 0.0345  
Epoch 6000, Loss: 0.0178  
Epoch 8000, Loss: 0.0115  
Epoch 10000, Loss: 0.0082

---

---

### Training with TANH

---

---

Epoch 2000, Loss: 0.0891  
Epoch 4000, Loss: 0.0241  
Epoch 6000, Loss: 0.0120  
Epoch 8000, Loss: 0.0075  
Epoch 10000, Loss: 0.0053

---

---

### Training with RELU

---

---

Epoch 2000, Loss: 0.0654  
Epoch 4000, Loss: 0.0198  
Epoch 6000, Loss: 0.0101  
Epoch 8000, Loss: 0.0065  
Epoch 10000, Loss: 0.0046

---

---

## PERFORMANCE COMPARISON

---

---

SIGMOID: Loss=0.0082, Accuracy=100%

Input	Target	Pred	Prob
[0 0]	0	0	0.042
[0 1]	1	1	0.963
[1 0]	1	1	0.962
[1 1]	0	0	0.039

TANH: Loss=0.0053, Accuracy=100%

Input	Target	Pred	Prob
[0 0]	0	0	0.031
[0 1]	1	1	0.971
[1 0]	1	1	0.971
[1 1]	0	0	0.030

RELU: Loss=0.0046, Accuracy=100%

Input	Target	Pred	Prob
[0 0]	0 0	0.027	
[0 1]	1 1	0.975	
[1 0]	1 1	0.975	
[1 1]	0 0	0.026	

## 4. OBSERVATIONS

### 1. Convergence Speed:

- ReLU: Fastest (Loss: 0.0046)
- Tanh: Moderate (Loss: 0.0053)
- Sigmoid: Slowest (Loss: 0.0082)

2. Accuracy: All achieved 100% on XOR problem

### 3. Gradient Flow:

- ReLU: No vanishing gradient problem
- Tanh: Better than Sigmoid, centered at 0
- Sigmoid: Slower due to gradient saturation

4. Decision Boundaries: All create non-linear boundaries separating XOR classes

---

## 5. CONCLUSION

Successfully implemented gradient descent and backpropagation with three activation functions. ReLU showed fastest convergence and best final loss, while all functions solved the non-linear XOR problem with 100% accuracy. The choice of activation function significantly impacts training speed and gradient flow.